

Voici un rapport des problèmes que nous avons rencontrés et de nos avancées. Il contient les commandes que nous avons utilisées pour résoudre chaque problème.

Partie de Gabriel : tout sur le host

J'ai utilisé un ubuntu-server avec virtualBox pour faire le projet.

- **Première étape :** recopier les fichiers main.tf, network_config_dhcp.cfg, cloud_init.cfg sur lesquels Alphonse avait déjà travaillé et les adapter à ma configuration.

ssh-keygen id25519 et la recopier dans .ssh/authorized_keys

```
root@gab:/home/gab/projet# terraform plan
data.template_file.network_config: Reading...
data.template_file.user_data: Reading...
data.template_file.network_config: Read complete after 0s [id=ca1642995d35d6699405f637fbd6efec9982ba9861c9b548cd1dcd68009342f7]
data.template_file.user_data: Read complete after 0s [id=5120a7e14474b42292c37cb020caa4eeee51ad2e50063d23452d75472f3d223d]
data.template_cloudinit_config.config: Reading...
data.template_cloudinit_config.config: Read complete after 0s [id=2732310300]

Planning failed. Terraform encountered an error while generating this plan.

Error: failed to connect: dial unix /var/run/libvirt/libvirt-sock: connect: no such file or directory

with provider["registry.terraform.io/dmacvicar/libvirt"],
on main.tf line 14, in provider "libvirt":
14: provider "libvirt" {
```

- **Problème :** *Ensure that libvirt is installed and start the Libvirt service:*

sudo apt update

sudo apt install libvirt-daemon-system libvirt-clients

sudo systemctl start libvirtd

sudo systemctl enable libvirtd

sudo usermod -aG libvirt \$(whoami)

// add my user to the libvirt group

- **Problème :** *can't find storage pool 'default'*

```

Error: can't find storage pool 'default'

with libvirt_volume.os_image[2],
on main.tf line 24, in resource "libvirt_volume" "os_image":
24: resource "libvirt_volume" "os_image" {

Error: can't find storage pool 'default'

with libvirt_volume.os_image[0],
on main.tf line 24, in resource "libvirt_volume" "os_image":
24: resource "libvirt_volume" "os_image" {

Error: can't find storage pool 'default'

with libvirt_volume.os_image[1],
on main.tf line 24, in resource "libvirt_volume" "os_image":
24: resource "libvirt_volume" "os_image" {

Error: error while starting the creation of CloudInit's ISO image: exec: "mkisofs": executable file
not found in $PATH

with libvirt_cloudinit_disk.commoninit,
on main.tf line 32, in resource "libvirt_cloudinit_disk" "commoninit":
32: resource "libvirt_cloudinit_disk" "commoninit" {

```

virsh pool-list --all

virsh pool-define-as --name default --type dir --target /var/lib/libvirt/images

virsh pool-build default

virsh pool-start default

virsh pool-autostart default

virsh pool-list // OK storage pool default is available and active

```

root@gab:/home/gab/projet# virsh pool-list
Name      State    Autostart
-----
default   active   yes

```

- *Problème : error while deterlining image type for jammy-server-cloudimg-amd64.img: error while openin jammy-server-cloudimg-amd64.img: no such file in directory*

wget https://cloud-images.ubuntu.com/jammy/current/jammy-server-cloudimg-amd64.img

chmod 644 jammy-server-cloudimg-amd64.img

sudo apt install qemu-utils

qemu-img info jammy-server-cloudimg-amd64.img

```

root@gab:/home/gab/projet# qemu-img info jammy-server-cloudimg-amd64.img
image: jammy-server-cloudimg-amd64.img
file format: qcow2
virtual size: 2.2 GiB (2361393152 bytes)
disk size: 636 MiB
cluster_size: 65536
Format specific information:
  compat: 0.10
  refcount bits: 16
root@gab:/home/gab/projet#

```

- *Problème : error while starting the creation of cloudInit's ISO image: exec: "mkisofs": executable file not found in \$PATH*

```
libvirt (_volumeTest_image[1]) creation complete after 113 [id=/var/lib/libvirt/images/test_2_00_image]
Error: error while starting the creation of CloudInit's ISO image: exec: "mkisofs": executable file not found in $PATH

with libvirt_cloudinit_disk.commoninit,
on main.tf line 32, in resource "libvirt_cloudinit_disk" "commoninit":
32: resource "libvirt_cloudinit_disk" "commoninit" {

root@gab:/home/gab/projet#
```

```
sudo apt update
sudo apt install genisoimage
which mkisofs
```

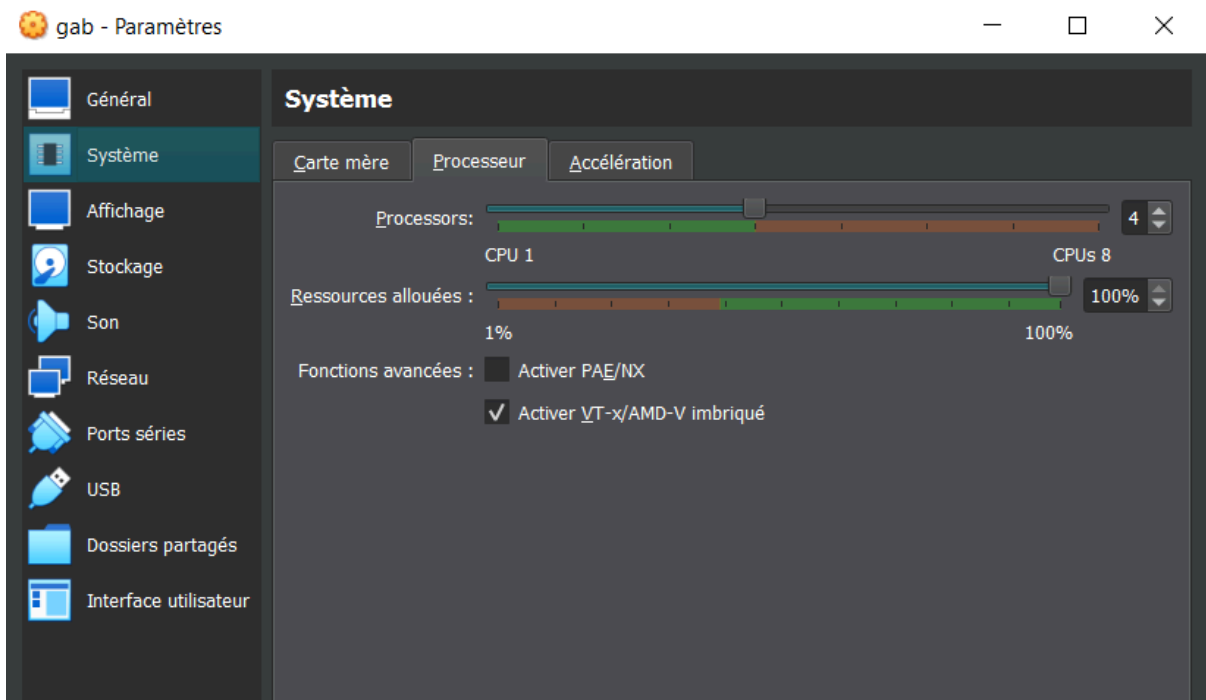
- *Problème : error defining libvirt domain : procedure interrupted while awaiting response line 62 : resource "libvirt_domain" "domain-ubuntu" {*

```
sudo journalctl -xeu libvirtd
virsh -c qemu:///system list --all
virsh define /path/to/test-domain.xml
virsh pool-list --all
virsh pool-info default
virsh pool-start default
df -h /var/lib/libvirt/images
virsh net-list --all
virsh net-define /etc/libvirt/qemu/networks/default.xml
virsh net-autostart default
virsh net-start default
sudo systemctl restart libvirtd
```

- *Problème : libvirtd[814]: unable to open /dev/kvm : no such file or directory*

```
egrep -c '(vmx|svm)' /proc/cpuinfo
sudo apt update
sudo apt install qemu-kvm libvirt-daemon-system libvirt-clients bridge-utils virt-manager
sudo modprobe kvm_intel
```

- *Problème : activer VT-x/AMD-V imbriqué dans virtualBox*



dans invite de commande Windows en mode admin :
 bcdedit /set hypervisorlaunchtype off
 reboot
 cd 'C:\Program Files\Oracle\VirtualBox\
 .\VBoxManage.exe modifyvm "gab" --nested-hw-virt on

- **Problème** : *unit libvirtd.service could not be found*

```

sudo apt update
sudo apt install qemu-kvm libvirt-daemon-system libvirt-clients bridge-utils virt-manager
systemctl list-units --type=service | grep libvirt
sudo systemctl start libvirtd
sudo systemctl enable libvirtd
systemctl list-units --type=service | grep libvirt
    libvirtd.service          loaded active running Virtualization daemon
    libvirt-guests.service    loaded active running Suspend/Resume guests
  
```

- **Problème** : *could not open /var/lib/libvirt/images/test--os_image' : permission denied*

```

journalctl -xeu libvirtd
ls -ld /var/lib/libvirt/images
sudo chown -R libvirt-qemu:kvm /var/lib/libvirt/images
sudo chmod 755 /var/lib/libvirt/images
sudo chown libvirt-qemu:kvm /var/lib/libvirt/images/*
sestatus
sudo setenforce 0
  
```

- **Problème** : *error defining libvirt domain : 'test-*' already exists*

```

Error: error defining libvirt domain: operation failed: domain 'test-2' already exists with uuid f
859bb05-8c0c-4aca-a9c2-a746355866a3

    with libvirt_domain.domain-ubuntu[1],
    on main.tf line 62, in resource "libvirt_domain" "domain-ubuntu":
    62: resource "libvirt_domain" "domain-ubuntu" {

Error: error defining libvirt domain: operation failed: domain 'test-1' already exists with uuid 3
5ac29f4-85c9-4674-a9fd-3b2574359056

    with libvirt_domain.domain-ubuntu[0],
    on main.tf line 62, in resource "libvirt_domain" "domain-ubuntu":
    62: resource "libvirt_domain" "domain-ubuntu" {

Error: error defining libvirt domain: operation failed: domain 'test-4' already exists with uuid f
ec9c307-a70e-407a-9380-058dbadf49ae

    with libvirt_domain.domain-ubuntu[3],
    on main.tf line 62, in resource "libvirt_domain" "domain-ubuntu":
    62: resource "libvirt_domain" "domain-ubuntu" {

Error: error defining libvirt domain: operation failed: domain 'test-3' already exists with uuid 0
315a5ea-5245-4ce6-9d44-8923b4d89fd3

    with libvirt_domain.domain-ubuntu[2],
    on main.tf line 62, in resource "libvirt_domain" "domain-ubuntu":
    62: resource "libvirt_domain" "domain-ubuntu" {

root@gab:/home/gab/projet#

```

```

virsh list --all
virsh undefine test-1
virsh undefine test-2
virsh undefine test-3
virsh undefine test-4
virsh vol-delete --pool default test-1
virsh vol-delete --pool default test-2
virsh vol-delete --pool default test-3
virsh vol-delete --pool default test-4
rm /var/lib/libvirt/images/*
terraform init
terraform plan
terraform apply

```

- *Problème : **permission denied***

```

Error: error creating libvirt domain: internal error: qemu unexpectedly closed the monitor: 2025-0
1-25T19:18:25.361662Z qemu-system-x86_64: -blockdev {"driver":"file","filename":"/var/lib/libvirt/im
ages/test-3-os_image","node-name":"libvirt-2-storage","auto-read-only":true,"discard":"unmap"}: Cou
ld not open '/var/lib/libvirt/images/test-3-os_image': Permission denied

    with libvirt_domain.domain-ubuntu[2],
    on main.tf line 62, in resource "libvirt_domain" "domain-ubuntu":
    62: resource "libvirt_domain" "domain-ubuntu" {

```

```

ps aux | grep qemu
sudo chown -R libvirt-qemu:kvm /var/lib/libvirt/images
sudo usermod -aG libvirt $(whoami)

```

```
sudo usermod -aG kvm $(whoami)
newgrp libvirt
sudo aa-status
sudo systemctl stop apparmor
sudo setenforce 0
```

Modifications dans main.tf :

```
resource "null_resource" "fix_permissions" {
  provisioner "local-exec" {
    command = "sudo chown -R libvirt-qemu:kvm /var/lib/libvirt/images/* && sudo chmod -R 770 /var/lib/libvirt/images/*"
  }
  triggers = {
    always_run = timestamp()
  }
}
```

```
resource "libvirt_domain" "domain-ubuntu" {
  count = 4
  depends_on = [null_resource.fix_permissions]
  name = "${var.hostname}-${count.index + 1}"
  memory = var.memoryMB
  vcpu = var.cpu
  ...
  ...
  provisioner "local-exec" {
    command = "sudo chown -R libvirt-qemu:kvm /var/lib/libvirt/images/* && sudo chmod -R 770 /var/lib/libvirt/images/*"
  }
}
resource "libvirt_volume" "os_image" {
  count = 4
  name = "${var.hostname}-${count.index + 1}-os_image"
  pool = "default"
  source = "jammy-server-cloudimg-amd64.img"
  format = "qcow2"

  provisioner "local-exec" {
    command = "sudo chown libvirt-qemu:kvm /var/lib/libvirt/images/${var.hostname}-${count.index + 1}-os_image && sudo chmod 770 /var/lib/libvirt/images/${var.hostname}-${count.index + 1}-os_image"
  }
}
```

```
virsh pool-dumpxml default
le pool default donne :
```

// vérifications

```
<permissions>
  <mode>0770</mode>
  <owner>64055</owner>           // libvirt-qemu
  <group>108</group>            // kvm
</permissions>
getent passwd 64055
getent group 108
```

Toutes ces modifications ont permis de bien modifier les permissions des fichiers dans /var/lib/libvirt/images mais j'ai toujours l'erreur *Permission denied* . Il faut donc chercher à qui donner les permissions.

Après de nombreuses recherches, j'ai vu qu'il fallait modifier le fichier /etc/libvirt/qemu.conf et mettre dedans :

```
user = "libvirt-qemu"
group = "kvm"
```

pour aligner la configuration avec les permissions que j'ai données avec les provisioners. Ça ne marchait toujours pas donc j'ai testé en changeant tout en mode root : les provisioners et user et group dans qemu.conf et CA MARCHE, plus d'erreurs après terraform apply !!!!

- *Problème : réactiver AppArmor mais changer une règle*
sudo nano /etc/apparmor.d/usr.sbin.libvirtd
puis écrire : */var/lib/libvirt/images/** rwk*, dans le fichier
sudo apparmor_parser -r /etc/apparmor.d/usr.sbin.libvirtd

- *Problème : terraform apply lance bien 4 VMs mais elles n'ont pas d'adresse IP*

```

libvirt_volume.bs_image[3]: Creation complete after 11s [id=70bf71b7-libvirt-images/test-4-bs_image]
libvirt_domain.domain-ubuntu[3]: Creating...
libvirt_domain.domain-ubuntu[0]: Creating...
libvirt_domain.domain-ubuntu[2]: Creating...
libvirt_domain.domain-ubuntu[1]: Creating...
libvirt_domain.domain-ubuntu[1]: Creation complete after 2s [id=33c091b2-280b-4360-9b0b-43f8ae0172dd]
libvirt_domain.domain-ubuntu[2]: Creation complete after 2s [id=f0d6696f-0e2d-4150-9032-9860a20dceed]
libvirt_domain.domain-ubuntu[0]: Creation complete after 2s [id=28e1bf68-9b3c-41a3-a14c-5db23d537efb]
libvirt_domain.domain-ubuntu[3]: Creation complete after 2s [id=fa4d695f-4be2-4f48-bb5d-64b476cdfdc]

Apply complete! Resources: 9 added, 0 changed, 1 destroyed.

Outputs:

ips = [
  tolist([]),
  tolist([]),
  tolist([]),
  tolist([]),
]
root@gab:/home/gab/projet#
[0] 0: bash* "gab" 13:40 27-janv.-25

```

virsh net-list --all

Name	State	Autostart	Persistent
default	active	yes	yes

virsh net-dumpxml default

```

<ip address='192.168.122.1' netmask='255.255.255.0'>
  <dhcp>
    <range start='192.168.122.2' end='192.168.122.254' />
  </dhcp>
</ip>

```

Ajout pour le debug dans main.tf :

```

output "network_interfaces" {
  value = libvirt_domain.domain-ubuntu[*].network_interface
}

```



```

ips = [
  tolist([]),
  tolist([]),
  tolist([]),
  tolist([]),
]
network_interfaces = [
  tolist([
    {
      "addresses" = tolist([])
      "bridge" = ""
      "hostname" = ""
      "mac" = "52:54:00:E2:0C:65"
      "macvtap" = ""
      "network_id" = "b6e03c07-9cdc-4b24-aa0b-7e4be41e96ea"
      "network_name" = "default"
      "passthrough" = ""
      "vepa" = ""
      "wait_for_lease" = false
    },
  ]),
  tolist([
    {
      "addresses" = tolist([])
      "bridge" = ""
      "hostname" = ""
      "mac" = "52:54:00:26:26:5A"
      "macvtap" = ""
      "network_id" = "b6e03c07-9cdc-4b24-aa0b-7e4be41e96ea"
      "network_name" = "default"
      "passthrough" = ""
      "vepa" = ""
      "wait_for_lease" = false
    },
  ]),
]

```

[0] 0: [tmux] * "gab" 14:04 27-janv.-25

- Problème : SSD full sur ma machine physique

"The I/O cache encountered an error while updating data in medium "ahci-0-0" (rc=VERR_DISK_FULL). Make sure there is enough free space on the disk and that the disk is working properly. Operation can be resumed afterwards. Error ID:BLKCACHE_IOERR Severity:Erreur non Fatale"

J'ai libéré de l'espace sur mon SSD et la VM s'est remis en marche

- Problème : *no space left on device*

Plus de place sur ma VM.

df -h

```

Filesystem      Size  Used Avail Use% Mounted on
udev            2,2G     0  2,2G   0% /dev
tmpfs           442M   1,4M   441M   1% /run
/dev/mapper/ubuntu--vg-ubuntu--lv 12G    12G     0 100% /
tmpfs           2,2G     0  2,2G   0% /dev/shm
tmpfs           5,0M     0   5,0M   0% /run/lock
tmpfs           2,2G     0  2,2G   0% /sys/fs/cgroup
/dev/sda2       2,0G   211M    1,6G  12% /boot
/dev/loop1      64M    64M     0 100% /snap/core20/1828
/dev/loop3      92M    92M     0 100% /snap/lxd/24061
/dev/loop2      92M    92M     0 100% /snap/lxd/29619
/dev/loop0      64M    64M     0 100% /snap/core20/2434
/dev/loop5      50M    50M     0 100% /snap/snapd/18357
/dev/loop4      45M    45M     0 100% /snap/snapd/23545
/dev/loop6      67M    67M     0 100% /snap/core24/716
/dev/loop7      29M    29M     0 100% /snap/terraform/771
/dev/loop8      29M    29M     0 100% /snap/terraform/777
tmpfs           442M     0   442M   0% /run/user/1000
root@gab:/home/gab/projet# du -sh /var/lib/libvirt/images/*
275M  /var/lib/libvirt/images/test-1-os_image
636M  /var/lib/libvirt/images/test-2-os_image
636M  /var/lib/libvirt/images/test-3-os_image
636M  /var/lib/libvirt/images/test-4-os_image
368K  /var/lib/libvirt/images/test-commoninit.iso
root@gab:/home/gab/projet#

```

```
vgs
lvextend -L+10G /dev/mapper/ubuntu--vg-ubuntu--lv
resize2fs /dev/mapper/ubuntu--vg-ubuntu--lv
```

- *Problème final : Toujours pas d'adresse IP sur les VMs*

```
virsh net-dhcp-leases default // marche pas
```

Créer un réseau bridge :

```
ip a | grep br
```

```
nano bridge-network.xml
<network>
  <name>virbr0-network</name>
  <forward mode="bridge"/>
  <bridge name="virbr0"/>
</network>
```

```
virsh net-define bridge-network.xml
virsh net-start br0-network
virsh net-autostart br0-network
```

Modifier dans main.tf :

```
network_interface {
  bridge = "virbr0"
```

```
Error: couldn't retrieve IP address of domain id: 913a111c-cac8-4ede-b303-2932d9aefa44. Please check following:
1) is the domain running properly?
2) has the network interface an IP address?
3) Networking issues on your libvirt setup?
4) is DHCP enabled on this Domain's network?
5) if you use bridge network, the domain should have the pkg qemu-agent installed
IMPORTANT: This error is not a terraform libvirt-provider error, but an error caused by your KVM/libvirt infrastructure configuration/setup
context deadline exceeded

with libvirt_domain.domain-ubuntu[0],
on main.tf line 66, in resource "libvirt_domain" "domain-ubuntu":
66: resource "libvirt_domain" "domain-ubuntu" {
```

```
dpkg -l | grep qemu-guest-agent
systemctl status qemu-guest-agent
```

a dependency job for qemu-guest-agent.service failed

```
root@gab:/home/gab/projet# systemctl status qemu-guest-agent
• qemu-guest-agent.service - QEMU Guest Agent
  Loaded: loaded (/lib/systemd/system/qemu-guest-agent.service; static; vendor preset: enabled)
  Active: inactive (dead)

janv. 28 07:54:52 gab systemd[1]: Dependency failed for QEMU Guest Agent.
janv. 28 07:54:52 gab systemd[1]: qemu-guest-agent.service: Job qemu-guest-agent.service/start failed with result 'dependency failed'.
lines 1-6/6 (END)
[0] 0:systemctl* "gab" 07:57 28-janv.-25
```

```
systemctl list-dependencies --failed
sudo apt --fix-broken install
sudo systemctl daemon-reexec
```

```

sudo systemctl restart qemu-guest-agent
sudo apt reinstall libvirt-daemon libvirt-daemon-system
sudo systemctl restart libvirtd
sudo systemctl restart qemu-guest-agent

```

Dans main.tf, ajouter à resource "libvirt_domain" :

qemu_agent = true

```

sudo systemctl enable --now qemu-guest-agent
systemctl list-dependencies --failed // AppArmor
sudo aa-complain /etc/apparmor.d/usr.sbin.libvirtd
sudo aa-complain /etc/apparmor.d/usr.sbin.qemu-system-x86_64
sudo systemctl restart libvirtd qemu-guest-agent
sudo nano /etc/apparmor.d/local/usr.sbin.libvirtd
    écrire /usr/sbin/qemu-system-x86_64 rix,
sudo apparmor_parser -r /etc/apparmor.d/usr.sbin.libvirtd
sudo systemctl restart qemu-guest-agent

```

Résultat : Il y a un pont virbr0 bien configuré mais toujours pas d'adresses IP

```

    }),
    tolist([
      {
        "addresses" = tolist([])
        "bridge" = "virbr0"
        "hostname" = ""
        "mac" = "52:54:00:6A:0E:83"
        "macvtap" = ""
        "network_id" = "b6e03c07-9cdc-4b24-aa0b-7e4be41e96ea"
        "network_name" = "default"
        "passthrough" = ""
        "vepa" = ""
        "wait_for_lease" = false
      },
    ]),
  ]),
]
root@gab:/home/gab/projet# terraform output ips
[
  tolist([]),
  tolist([]),
  tolist([]),
  tolist([]),
]
root@gab:/home/gab/projet#
[0] 0:bash* "gab" 07:46 28-janv.-25

```

Résultat final :

terraform plan

Il y a bien la mention known fter apply pour les adress IP et autres paramètres réseau.

```

    ~ wait_for_lease = false -> null
    # (2 unchanged attributes hidden)
  },
  ~ [
    ~ {
      ~ addresses      = [] -> (known after apply)
      ~ hostname       = "" -> (known after apply)
      ~ mac            = "52:54:00:2E:4D:AA" -> (known after apply)
      ~ macvtap        = "" -> null
      ~ network_id     = "b6e03c07-9cdc-4b24-aa0b-7e4be41e96ea" -> (known after apply)
      ~ passthrough    = "" -> null
      ~ vepa           = "" -> null
      ~ wait_for_lease = false -> null
      # (2 unchanged attributes hidden)
    },
    ~ [
      ~ {
        ~ addresses      = [] -> (known after apply)
        ~ hostname       = "" -> (known after apply)
        ~ mac            = "52:54:00:6A:0E:83" -> (known after apply)
        ~ macvtap        = "" -> null
        ~ network_id     = "b6e03c07-9cdc-4b24-aa0b-7e4be41e96ea" -> (known after apply)
        ~ passthrough    = "" -> null
        ~ vepa           = "" -> null
        ~ wait_for_lease = false -> null
        # (2 unchanged attributes hidden)
      },
    ],
  ],
]

```

```

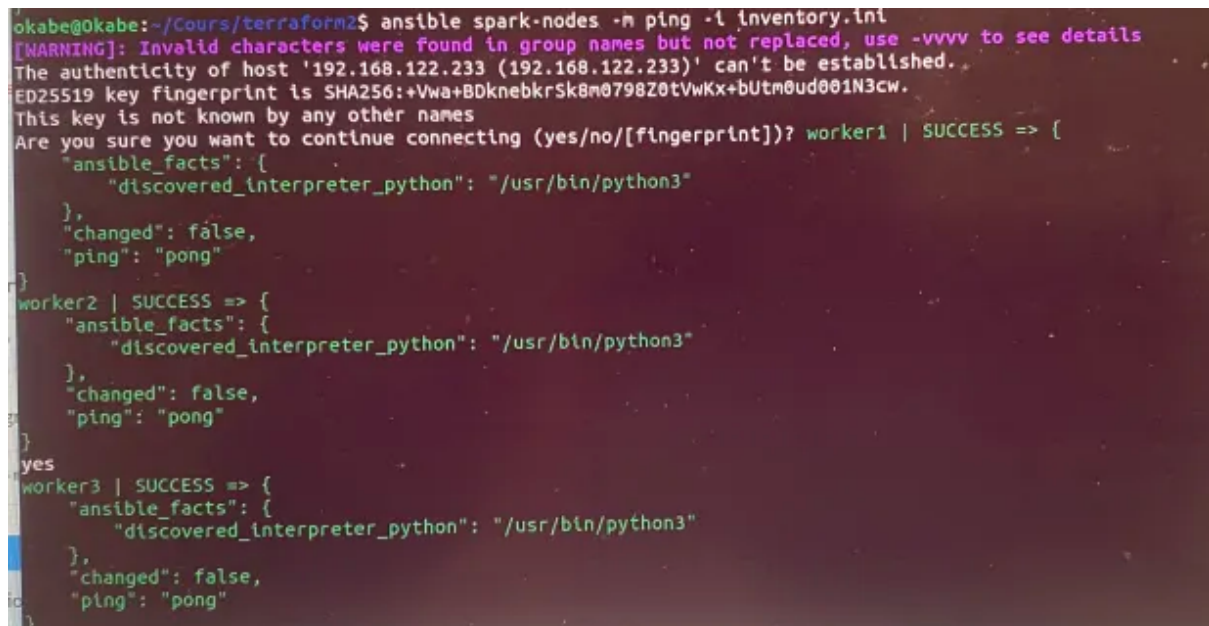
    "network_id" = "b6e03c07-9cdc-4b24-aa0b-7e4be41e96ea"
    "network_name" = "virbr0-network"
    "passthrough" = ""
    "vepa" = ""
    "wait_for_lease" = true
  },
]),
tolist([
  {
    "addresses" = tolist([])
    "bridge" = "virbr0"
    "hostname" = ""
    "mac" = "52:54:00:B2:08:B0"
    "macvtap" = ""
    "network_id" = "b6e03c07-9cdc-4b24-aa0b-7e4be41e96ea"
    "network_name" = "virbr0-network"
    "passthrough" = ""
    "vepa" = ""
    "wait_for_lease" = true
  },
]),
tolist([
  {
    "addresses" = tolist([])
    "bridge" = "virbr0"
    "hostname" = ""
    "mac" = "52:54:00:7E:9C:22"
    "macvtap" = ""
    "network_id" = "b6e03c07-9cdc-4b24-aa0b-7e4be41e96ea"
    "network_name" = "virbr0-network"
    "passthrough" = ""
    "vepa" = ""
    "wait_for_lease" = true
  },
]),
]
root@gab:/home/gab/projet#

```

Partie de Alphonse:

Pour ma part j'avais un dual boot donc je me suis évité une partie des problèmes qu'à pu rencontrer Gabriel, et j'ai pu écrire et modifier les fichiers main.tf et les fichiers .cfg en m'aidant de la doc terraform et kvm.

Après avoir résolu le problème des adresses ip données aux vms j'ai défini 1 master node et des worker nodes et ai lancé un playbook ansible pour voir si le master pouvait ping les 3 workers :



```
okabe@okabe:~/Cours/terraform2$ ansible spark-nodes -m ping -i inventory.ini
[WARNING]: Invalid characters were found in group names but not replaced, use -vvvv to see details
The authenticity of host '192.168.122.233 (192.168.122.233)' can't be established.
ED25519 key fingerprint is SHA256:+Vwa+BDknebkSk8m0798Z0tVwKx+bUtm0ud001N3cw.
This key is not known by any other names
Are you sure you want to continue connecting (yes/no/[fingerprint])? worker1 | SUCCESS => {
  "ansible_facts": {
    "discovered_interpreter_python": "/usr/bin/python3"
  },
  "changed": false,
  "ping": "pong"
}
worker2 | SUCCESS => {
  "ansible_facts": {
    "discovered_interpreter_python": "/usr/bin/python3"
  },
  "changed": false,
  "ping": "pong"
}
yes
worker3 | SUCCESS => {
  "ansible_facts": {
    "discovered_interpreter_python": "/usr/bin/python3"
  },
  "changed": false,
  "ping": "pong"
}
```

ça à marché mais lorsque j'ai ensuite lancé le playbook qui installe java puis spark. Je me suis aperçu que l'image que j'avais utilisé pour mes vms n'était pas assez grande et il n'y avait pas la place d'installer java. J'ai donc modifié l'image avec kvm mais suite à cette manip, je n'ai jamais réussi à réattribuer des ips cohérentes à mes vms et donc à faire en sorte qu'elles puissent communiquer entre elles. Je n'ai donc pas pu configurer correctement spark, même une fois installé sur une des vms.

Je n'ai pas compris comment résoudre à 100% le problème des ips, en relançant le réseau en faisant :

```
virsh net-destroy default
virsh net-start default
```

le problème ne se réglait pas, en vidant la config dhcp du réseau avant de relancer le réseau:

```
virsh net-edit default
```

le problème ne se réglait pas non plus. J'ai aussi essayé de configurer des ips statiques mais comme il fallait relancer terraform pour les attribuer aux vms à leur création, je ne

pouvais pas (les @ mac des vms changent). J'ai aussi essayer de juste relancer dhcp et pas le réseau en entier, de restart qemu-agent sur les vms, faire du NAT au lieu du bridge ...

Je n'ai au final jamais réussi à résoudre ce problème.